


Introducing Switch: a framework for custom data applications

Josh Ferguson
Chief Architect @ Mode



**we're going to talk about building tools
to make better decisions with data**

actually we're going to talk about **building
tools to build tools** to make better decisions
with data



**i've been obsessed with building these
kinds of tools for about 20 years**

@besquared almost everywhere



mode(.com)



a collaborative data science platform

**our users are data scientists, analysts, and
engineers**

**whose job it is to help everybody make better
decisions with data**

we make visualization tools

**a few years ago we decided to rethink our
approach to building these tools**

from static to explorable

from presentation to analysis

we had to build a better data framework

we call that framework switch

**it's a collection of typescript libraries and
tools that accelerate **our** data
application development**

**we're in the process of releasing the
framework with an open source license**

**we think that everybody should
build data applications with it**

what's a data application?

a customer data portal

an a/b testing tool for our product team

an account lookup app for sales

**a customer engagement scoring tool for
success**


a revenue modeling dashboard for finance



why would we build these ourselves?

can't we just buy them?

**Well supported by
companies and tools**



**there's no collection of off-the-shelf tools
that will provide everything we need to
make better decisions with data**

**as our organizations continue to become
more analytically mature, we need to build
custom data apps to keep up**

sounds great, what's the problem?

today almost all of these apps are one-offs



how can we make our data apps faster to build, easier to maintain, and more reliable?

**there's several challenges we've got to
address to get there**

we're going to look at three of them

how the framework addresses each one

challenge number one

CHALLENGE #1

our data is incomplete and messy

CHALLENGE #1


we don't know what we'll need ahead of time

CHALLENGE #1

we can't build a new etl pipeline or deploy our app **every time** we need to answer a slightly different question

CHALLENGE #1

**we should let our users quickly and easily
express data in new and different ways**



Introducing Formulas

an excel-like language for data expression

**they lets us work with datasets, even if we're not
database or programming language experts**

what can we do with them?

**unlike excel whose formulas operate on cells, our
formulas operate on entire datasets at a time**

FORMULAS

sample dataset

ID	Date	Product	Quantity	Price	Filled
1	2019-01-01	A	10	10.00	true
2	2019-01-02	B	5	20.00	false
...	

EXAMPLES

calculate ratios!

`[Price] / [Quantity]`

EXAMPLES

convert units!

Dollar to cents

`[Price] * 100`

EXAMPLES

clean data!

```
CASE [Product]
WHEN "A,"
THEN "A"
ELSE [Product]
END
```

EXAMPLES

aggregate data!

```
AVG([Price] / [Quantity])
```

EXAMPLES

lookup values!

```
LOOKUP(AVG([Price]), FIRST())
```

what else!?

let's look at the language to find out

LITERALS

nulls

NULL

LITERALS

booleans

TRUE

FALSE

LITERALS

numbers

-42

1000

3.1415926

0xBEEF

LITERALS

strings

`'Category'`

`"Product Name"`

LITERALS

dates

`#2019-04-18#`

`#2019-04-18T10:50:15#`

LITERALS

regular expressions

```
/[\w\d]+/ig
```

ACCESS

data access

[Product]
[Quantity]

OPERATORS

mathematic

```
[Quantity] * 500  
[Quantity] / 500  
[Quantity] + 500  
[Quantity] - 500  
[Quantity] % 500
```


OPERATORS

relational

[Quantity] = 500

[Quantity] <> 500

[Quantity] < 500

[Quantity] <= 500

[Quantity] > 500

[Quantity] >= 500

OPERATORS

logical

```
NOT [Filled]
```

```
[Filled] AND [Quantity] > 500
```

```
[Filled] OR [Quantity] <= 500
```

CONDITIONAL

if/elsif/else

```
IF [Quantity] > 500  
  THEN "Large Order"  
ELSIF [Quantity] > 250  
  THEN "Medium Order"  
ELSE "Small Order"  
END
```

CONDITIONAL

case

```
CASE [Filled]
WHEN TRUE
THEN "Filled"
WHEN FALSE
THEN "Unfilled"
ELSE "Unknown"
END
```

FUNCTIONS

constant

NOW()

FUNCTIONS

scalar

`FLOOR([Price])`

`TRIM([Product])`

`DATETRUNC('day', [Date])`

FUNCTIONS

aggregate

`SUM([Price])`

`AVG([Quantity])`

`COUNTD([Product])`

FUNCTIONS

analytic

```
RANK(SUM([Quantity]))
```

```
RUNNING_SUM(COUNT([Price]))
```

```
LOOKUP(AVG([Price]), FIRST())
```


that's it, simple and powerful

the possibilities are infinite

**formulas play a key role in the framework
and in our data applications**

**in the framework we use them to model,
query, and process our data**

they provide a **common language** that we
can use across all of our data applications

**we can build interfaces that let users
extend our apps with their own business
logic and calculations**

a single formula that takes someone a few
minutes to write might take hours or days
to implement and deploy otherwise

not having to build ETL pipelines or write
app code all the time can amplify our
effort **100x**

that's pretty rad

**let's keep going and see how we use formulas
to query our data**

challenge number two

CHALLENGE #2

getting from data to visualization

CHALLENGE #2

**we can't build ad-hoc data transformation code
every time we want to build a new visualization**

CHALLENGE #2

**we should describe the data we need in way
that matches the visualizations we're trying
to build**



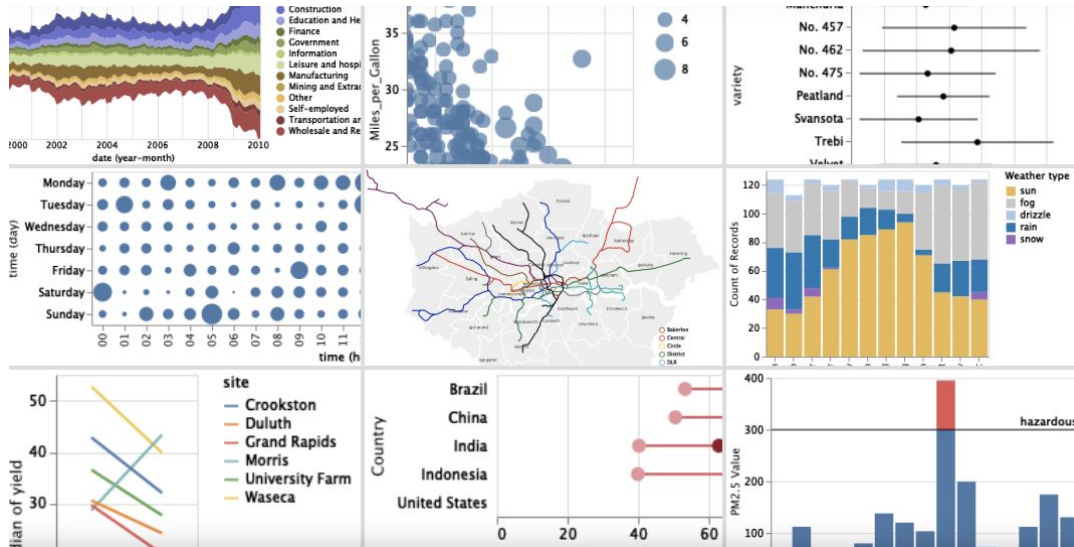
Introducing Queries

**our queries speak the language of
data visualization**

what language is that?

QUERIES

none other than
the grammar
of graphics



most of the visualizations that we can
encode with tools like vega-lite can be
translated directly into switch queries

how do they work?

first we define the data we want to see

QUERIES

**we start with fields
which are defined
with a formula**

```
Field {  
    formula: string;  
}
```

QUERIES

they let us
describe the data
and calculations
we want to see in
our result

```
Field {  
    formula: string;  
}
```

QUERIES

there are two
special fields
worth mentioning
which are called
names and values

```
Names {  
  formula: "$[Names]";  
}
```

```
Values {  
  formula: "$[Values]";  
}
```


NAMES/VALUES

they let us combine
multiple aggregate
fields together into
a single field

```
Names {  
  formula: "$[Names]";  
}  
  
Values {  
  formula: "$[Values]";  
}
```

NAMES/VALUES

we'll go through
an example of
where that can be
useful in a minute

```
Names {  
  formula: "$[Names]";  
}  
  
Values {  
  formula: "$[Values]";  
}
```

QUERIES

next we've got
filters

```
Filter {  
  field: Field;  
  conds: Conditions;  
}
```

QUERIES

they let us get rid
of data we don't
want in our result
by adding
conditions on our
fields

```
Filter {  
  field: Field;  
  conds: Conditions;  
}
```

QUERIES

last but not least
we have sorts

```
Sort {  
  field: Field;  
  type: SortType;  
  order: SortOrder;  
}
```

QUERIES

they let us
re-arrange our
result by adding
orders to our fields

```
Sort {  
  field: Field;  
  type: SortType;  
  order: SortOrder;  
}
```

then we map our data to our visualization

QUERIES

the first way to do
that is with marks

```
Mark {  
  field: Field;  
  color: Field[];  
  size: Field[];  
  label: Field[];  
  ...  
}
```


QUERIES

**marks are how we
define the layers
of our visualization**

```
Mark {  
  field: Field;  
  color: Field[];  
  size: Field[];  
  label: Field[];  
  ...  
}
```

QUERIES

1 mark = 1 layer

```
Mark {  
    field: Field;  
    color: Field[];  
    size: Field[];  
    label: Field[];  
    ...  
}
```

MARKS

it's defined by a
single field

```
Mark {  
  field: Field;  
  color: Field[];  
  size: Field[];  
  label: Field[];  
  ...  
}
```

MARKS

it's got channels
like color, size,
and label, that let
us associate
data with
visual properties

```
Mark {  
  field: Field;  
  color: Field[];  
  size: Field[];  
  label: Field[];  
  ...  
}
```

MARKS

we can map as many channels as we want based on the needs of our visualization

```
Mark {  
  field: Field;  
  color: Field[];  
  size: Field[];  
  label: Field[];  
  ...  
}
```

QUERIES

using marks and
the other pieces
we talked about
we can build a
complete visual
mapping which we
call a pivot query

```
PivotQuery {  
  column: Field[];  
    x: Field[];  
  row: Field[];  
    y: Field[];  
  
  values: Field[];  
  
  marks: Mark[];  
  filters: Filter[];  
  sorts: Sort[];  
}
```

QUERIES

we've got our
marks, filters, and
sorts

```
PivotQuery {  
  column: Field[];  
    x: Field[];  
  row: Field[];  
    y: Field[];  
  
  values: Field[];  
  
  marks: Mark[];  
  filters: Filter[];  
  sorts: Sort[];  
}
```

PIVOT QUERY

we've also got
some other
channels here that
we can use to map
other visual
properties

```
PivotQuery {  
  column: Field[];  
    x: Field[];  
  row: Field[];  
    y: Field[];  
  
  values: Field[];  
  
  marks: Mark[];  
  filters: Filter[];  
  sorts: Sort[];  
}
```


PIVOT QUERY

first, column and row which let us facet or group data across or down our visualization respectively

```
PivotQuery {  
  column: Field[];  
    x: Field[];  
  row: Field[];  
    y: Field[];  
  
  values: Field[];  
  
  marks: Mark[];  
  filters: Filter[];  
  sorts: Sort[];  
}
```

PIVOT QUERY

next, x and y
which let us
position data
across or down
our visualization
within those facets

```
PivotQuery {  
  column: Field[];  
    x: Field[];  
  row: Field[];  
    y: Field[];  
  
  values: Field[];  
  
  marks: Mark[];  
  filters: Filter[];  
  sorts: Sort[];  
}
```

PIVOT QUERY

finally, we've got values which let's us combine all of the fields in it into a single field that we can use in the other channels


```
PivotQuery {  
  column: Field[];  
    x: Field[];  
  row: Field[];  
    y: Field[];  
  
  values: Field[];  
  
  marks: Mark[];  
  filters: Filter[];  
  sorts: Sort[];  
}
```

Let's visualize an example

QUERIES

here's a
beautiful chart

Orders



QUERIES

a beautiful chart
deserves a
beautiful query

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXAMPLE

we've got day
on the x axis

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXAMPLE

our values field
on the y axis

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```


EXAMPLE

there's the sum of
price and a
running sum of
quantity in values

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXAMPLE

this chart is going
to be a single
layer so we've got
one mark

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXAMPLE

it's defined by
our values field


```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXAMPLE

within that layer we want to see two distinct series each with its own color so we add names to our color channel

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

Orders



```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

over time it becomes second nature

**once we learn to speak the language our
ability to quickly transform and visualize
data is increased by 10x**

challenge number three

CHALLENGE #3


**our datasets are millions and billions of rows
and growing**

CHALLENGE #3

we can't constantly move it around or materialize every view we might need to analyze ahead of time

CHALLENGE #3

**we should work with data it as it exists in
the places where it already lives**



Introducing Processors

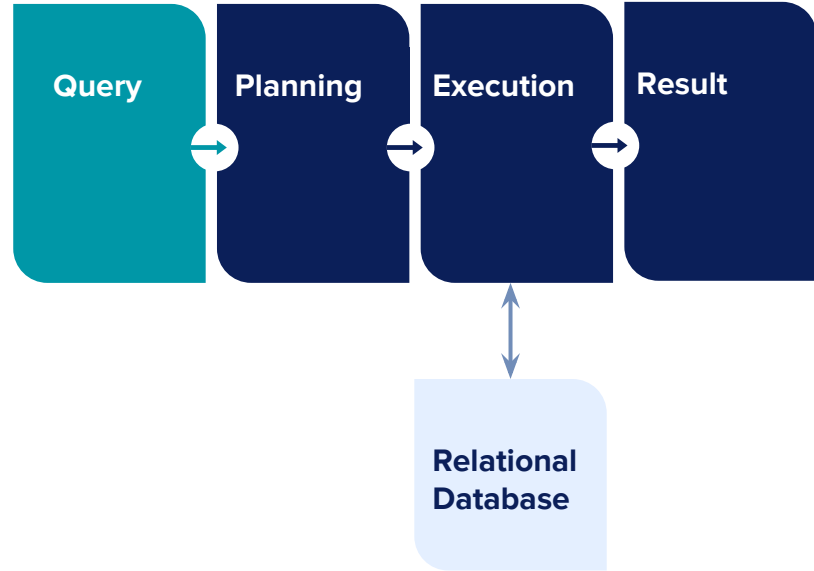
they're our database's analytical co-pilots

what do we mean by that?

let's talk about how they work

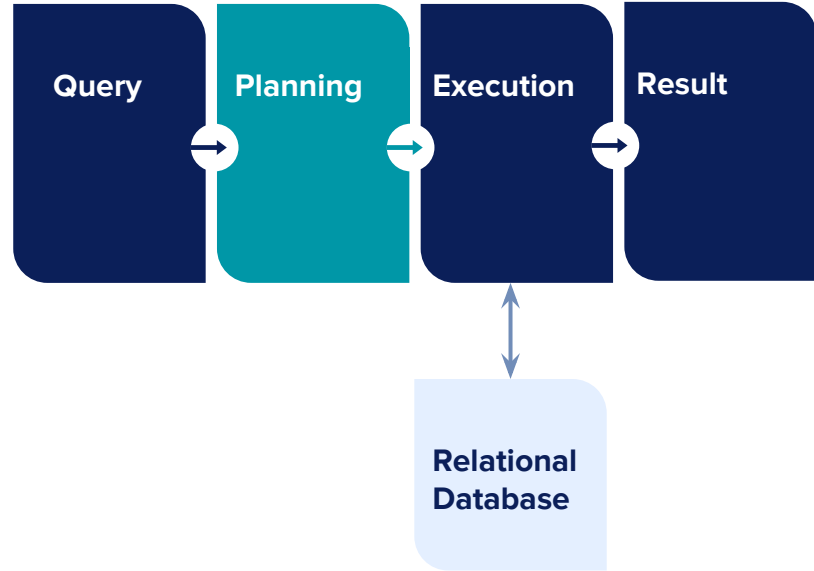
PROCESSORS

we start with a query like the one we saw in the last section




PROCESSORS

we build a
plan, which is
a set of
instructions
for processing
that query



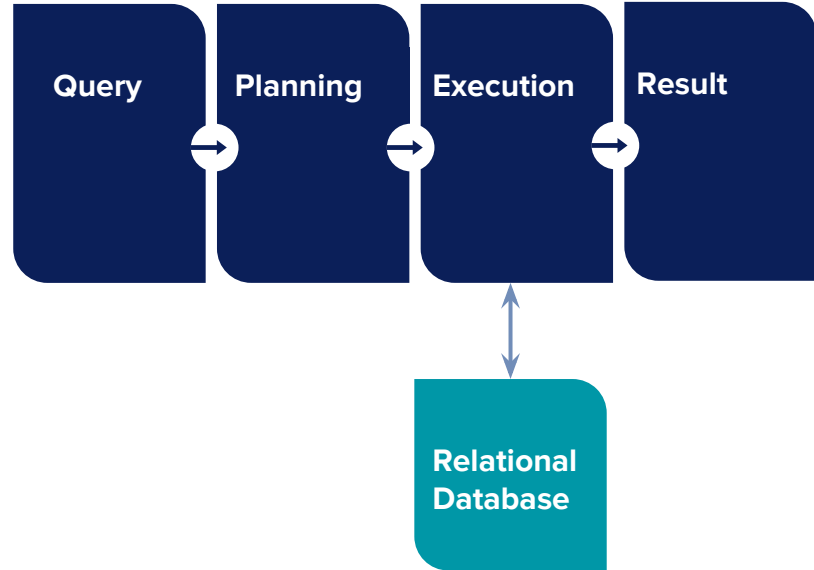
PROCESSORS

that plan gets passed along to the next step where it's executed by the processor



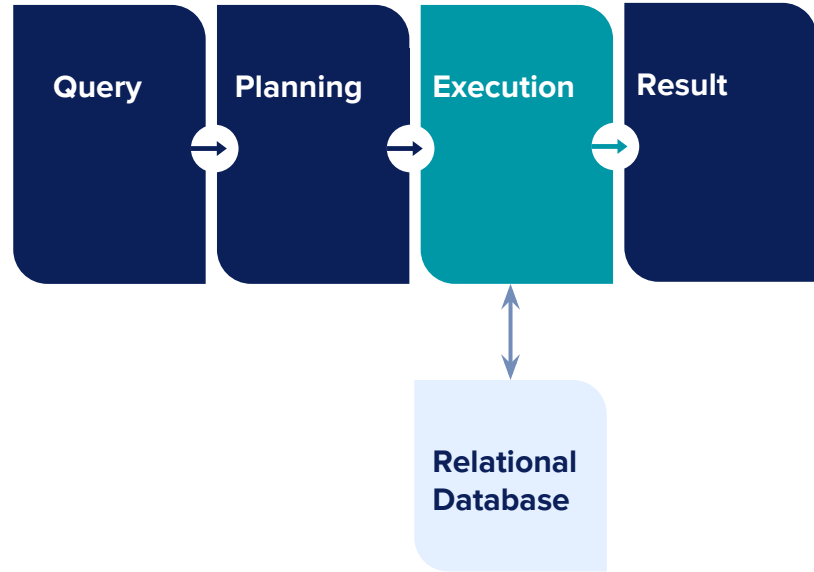
PROCESSORS

during
execution, the
processor will
issue queries
against our
database



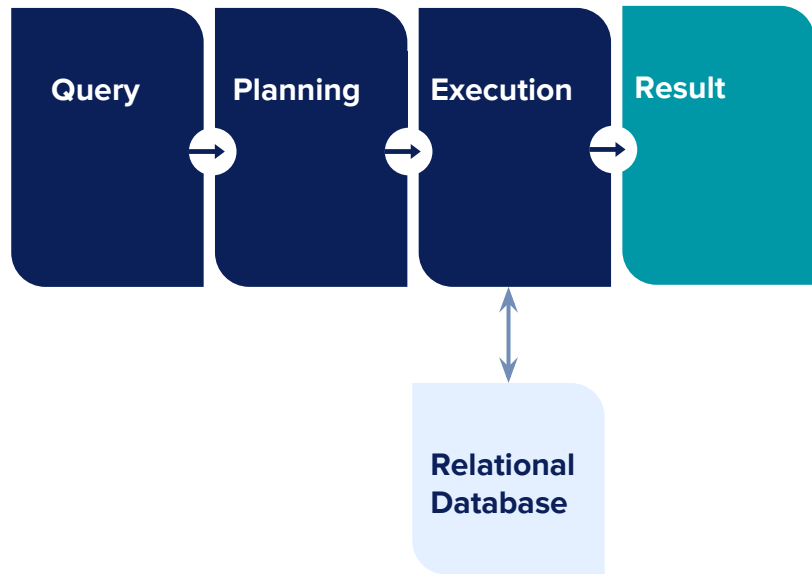
PROCESSORS

it'll take those intermediate query results and process them further to produce a final result



PROCESSORS

**the last step is
taking the final
result and
sending it back
to our app**




let's look at how planning works first

the planner looks at our query in a specific order and builds a logical execution plan


PROCESSORS

**first, we go
through the
fields in each
channel**




PROCESSORS

**next, if we
have names or
values fields
we add
instructions for
those**




PROCESSORS

after that we
plan all of the
filters




PROCESSORS

**followed by
sorts**



PROCESSORS

finally we add
a limit or offset
if they're part
of the query



**as we go through each step the planner
decides what we want to do in the database
and what we want to process ourselves**

how does it decide?

the planner always decides to “**push-down**”
grouping and aggregate expressions and
“**pull-up**” analytical expressions

an example is in our plans

DATAFLOW

let's say we've got
this field in our
query

Field

```
1 + RUNNING_AVG(SUM([Price]) + 1)
```

DATAFLOW

this fragment is an
aggregate
expression

Field

```
1 + RUNNING_AVG(SUM([Price]) + 1)
```

|-----|

DATAFLOW

an aggregate expression is any aggregate function and the operators attached to it

Field

```
1 + RUNNING_AVG(SUM([Price]) + 1)
```

|-----|

DATAFLOW

this fragment
is an analytic
expression

Field

1 + RUNNING_AVG(SUM([Price]) + 1)

|-----|

DATAFLOW

**an analytic
expression is any
analytic function
and the operators
attached to it**

Field

```
1 + RUNNING_AVG(SUM([Price]) + 1)  
|-----|
```

DATAFLOW

the planner will
split this field into
two parts

Field

`1 + RUNNING_AVG(SUM([Price]) + 1)`

Push-Down

?

Pull-Up

?

DATAFLOW

the first part is the
aggregate
expression that
gets pushed down
to the database

Field

`1 + RUNNING_AVG(SUM([Price]) + 1)`

Push-Down

`SUM([Price]) + 1 AS C1`

Pull-Up

?

DATAFLOW

the second part is
the analytic
expression that
gets pulled up to
the processor

Field

`1 + RUNNING_AVG(SUM([Price]) + 1)`

Push-Down

`SUM([Price]) + 1 AS C1`

Pull-Up

`1 + RUNNING_SUM([C1])`

DATAFLOW

expressions that
get pushed down
are given a unique
alias that we use
when we do our
analytical
processing

Field

```
1 + RUNNING_AVG(SUM([Price]) + 1)
```

Push-Down

```
SUM([Price]) + 1 AS c1
```

Pull-Up

```
1 + RUNNING_SUM([c1])
```

**this is what we mean when we say processors
are your database's analytical co-pilot**

I know what you're thinking

seems like a lot of trouble

why don't we everything in the database?

**organizations operate dozens of
databases across almost as many vendors**

**they don't all support the same features
in the same way**

**we want our processors to work the same
way everywhere**

**we want a common data processing model we
can rely on across all of the apps in our
organization**

we've taken a “lowest common denominator”
approach to solving this problem

as long as our databases can do the basic stuff like select, group, aggregate, filter, and sort data, we can handle the rest

**it's also a great opportunity to extend
our processors to do more with our data**

things like regression, clustering, anova, etc.

**this means delivering better applications more
quickly to more people**

time to execute an example


EXECUTION

we're going to
walk through how
we would execute
a plan for our
beautiful query

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```


PROCESSORS

**the first step is
to execute our
pushed down
query against
our relational
database**



EXECUTION

we've got three
expressions here
that get pushed
down

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXECUTION

the first is the
truncated date on
our x-axis

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXECUTION

the second is the
sum of price in
values

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXECUTION

the third is the
sum of quantity
used in the
running_sum
function also in
values

```
PivotQuery {  
  x: [ "DATETRUNC('day', [Date])" ],  
  y: [ "$[Values]" ],  
  values: [  
    "SUM([Price])",  
    "RUNNING_SUM(SUM([Quantity]))"  
  ],  
  marks: [{  
    field: "$[Values]",  
    color: [ "$[Names]" ]  
  }]  
}
```

EXECUTION

just like before,
these expressions
get pushed down
to the database
resulting in this
beautiful sql
query

```
SELECT DATETRUNC('day', date) AS C1  
       SUM(price) AS C2,  
       SUM(quantity) AS C3  
FROM orders  
GROUP BY DATETRUNC('day', date)
```

EXECUTION

the table name
comes from our
data model which is
used during
planning to tell the
processor about our
database schema

```
SELECT DATETRUNC('day', date) AS C1  
       SUM(price) AS C2,  
       SUM(quantity) AS C3  
FROM orders  
GROUP BY DATETRUNC('day', date)
```

EXECUTION

we've saved
details about that
part for next time

```
SELECT DATETRUNC('day', date) AS C1  
       SUM(price) AS C2,  
       SUM(quantity) AS C3  
FROM orders  
GROUP BY DATETRUNC('day', date)
```



EXECUTION

this is what our
database gives
us back

DAY(Date)	SUM(price)	SUM(quantity)
2019-01-01	10	15
2019-01-02	5	5
...

PROCESSORS

**we take that
and process our
analytic
expressions**




EXECUTION

this gives us a
new attribute as
we can see here

DAY(Date)	SUM(price)	SUM(quantity)	RUNNING_SUM(quantity)
2019-01-01	10	15	15
2019-01-02	5	5	20
...	

PROCESSORS

since we have
names and
values, we use
a fold transform
to “unpivot” the
result




EXECUTION

this adds two new fields, names and values, and also expands the number of rows in the result

Names	Values	DAY (Date)	SUM (price)	SUM (quantity)	RUNNING_SUM (quantity)
SUM(price)	10	2019-01-01	10	15	15
RUNNING_SUM (quantity)	15	2019-01-01	10	5	15
SUM(price)	5	2019-01-02	5	15	20
RUNNING_SUM (quantity)	20	2019-01-02	5	5	20
...

PROCESSORS

almost done!
next we slice
out any
intermediate
fields that we
don't want
which we call
selection




EXECUTION

that gives us
this beautiful
result

Names	Values	DAY(Date)
SUM(price)	10	2019-01-01
RUNNING_SUM(quantity)	15	2019-01-01
SUM(price)	5	2019-01-02
RUNNING_SUM(quantity)	20	2019-01-02
...

PROCESSORS

**last but not
least, we take
our result and
send it back to
the app**




```

PivotQuery {
  x: [ "DATETRUNC('day', [Date])" ],
  y: [ "$[Values]" ],
  values: [
    "SUM([Price])",
    "RUNNING_SUM(SUM([Quantity]))"
  ],
  marks: [{
    field: "$[Values]",
    color: [ "$[Names]" ]
  }]
}

```

Names	Values	DAY(Date)
SUM(price)	10	2019-01-01
RUNNING_SUM(quantity)	15	2019-01-01
SUM(price)	5	2019-01-02
RUNNING_SUM(quantity)	20	2019-01-02
...

and that's how the tables turn

not having to move data around or materialize
all our views ahead of time lets us effectively
use **1000x** more data

where does that leave us?

**using formulas is 100x faster than writing
etl pipelines or app code**

using queries is 10x faster than building ad-hoc transformation code for visualizations

**using processors lets us work with 1000x
more data**

we did the math and that's 1,000,000x

I rest my case, your honor

**in all seriousness we think this can be
incredibly valuable for data teams
everywhere**

where are we going to go from here?

ROADMAP

**where are
we going
from here?**

- **Release the code under open license**

ROADMAP

**where are
we going
from here?**

- Release the code under open license
- **Expand the built-in function library**

ROADMAP

**where are
we going
from here?**

- Release the code under open license
- Expand the built-in function library
- **Build out more real-world examples**

ROADMAP

**where are
we going
from here?**

- Release the code under open license
- Expand the built-in function library
- Built out more real-world examples
- **Expand our database adapter library**

ROADMAP

**where are
we going
from here?**

- Release the code under open license
- Expand the built-in function library
- Build out more real-world examples
- Expand our database adapter library
- **Integrate with open tools like DBT**

ROADMAP

**where are
we going
from here?**

- Release the code under open license
- Expand the built-in function library
- Build out more real-world examples
- Expand our database adapter library
- Integrate with open tools like DBT
- **Integrate with libraries like vega-lite**

ROADMAP

**where are
we going
from here?**

- Release the code under open license
- Expand the built-in function library
- Build out more real-world examples
- Expand our database adapter library
- Integrate with open tools like DBT
- Integrate with libraries like vega-lite
- **Build common components for frameworks like angular, react, native, etc.**

how do we get involved?

COMMUNITY

first of all, go here 

github.com/switch-data/community

COMMUNITY

AND HIT THE STAR BUTTON 

github.com/switch-data/community

COMMUNITY

you won't regret it
github.com/switch-data/community

COMMUNITY

**we're gonna update it with links, resources,
issues, slack chat info, etc.**

COMMUNITY

how can I get involved?

COMMUNITY

I want to go to meetups

COMMUNITY

I want to visit organizations

COMMUNITY

I want to understand our challenges



thank you again

Q&A

Come see me during office hours!

